January 26, 2026

The Agentic AI Inflection Point in Banking

An Operating Model Divide for 2026

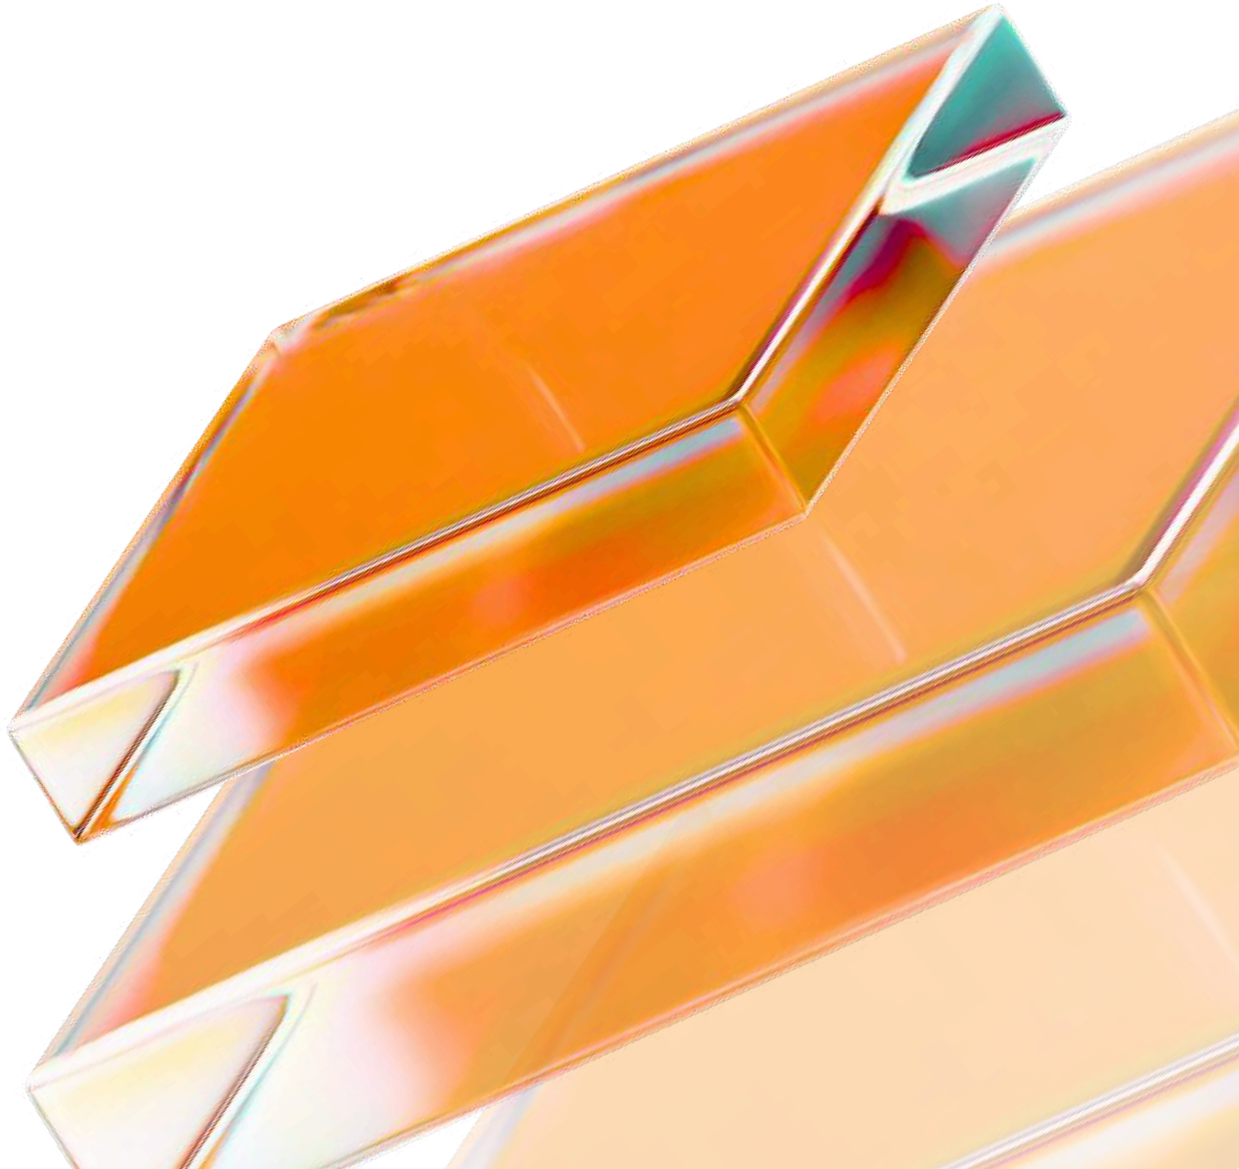

www.oticgroup.net

Executive Summary

Banks are now crossing a structural boundary: from **AI as an assistant** that answers questions and drafts content, to **AI as an autonomous actor** that plans, decides, and executes multi-step workflows within explicit risk, policy, and capital constraints. In this model, agents do not merely "support" staff; they own well-bounded production processes, invoke tools (core systems, risk engines, APIs), and escalate only true exceptions.

The central thesis is clear: **Agentic AI allows revenue and risk-adjusted volume growth to decouple from headcount growth for the first time in modern banking.** Early adopters are already seeing 20–40% structural reductions in middle- and back-office cost for agentified processes, 10–20% uplift in risk-adjusted revenue in selected businesses, and measurable reductions in operational losses and regulatory findings where agents run continuous controls rather than periodic checks. These gains are not coming from more chatbots; they are coming from autonomous remediation in credit, multi-step KYC resolution, intraday liquidity optimization, and self-healing IT and operations. [1,5,9,14]

By January 2026, three conclusions are unavoidable. First, **the experimentation phase is over:** table-stakes GenAI assistants are commoditized, and boards now demand efficiency ratio, ROE, and loss metrics. Second, **agentic systems are already in production** at leading global banks and fintechs, and they are compounding advantage. Third, **most institutions are structurally unprepared** to deploy agentic AI at scale given their legacy cores, fragmented data, and lack of orchestration and governance capabilities. Under these conditions, attempting to build end-to-end agentic capability internally is likely to fail on time-to-value, architecture, and talent. **Strategic partnerships with specialized transformation firms are emerging as the lowest-risk, highest-ROI path to close the capability chasm within the 24–36 month window shareholders and regulators will tolerate.**

Contents

1

The 2026 Agentic Landscape: Where Value Is Already Being Captured

From Automation to Orchestration: Defining "Agentic" in Banking Terms

Most banks today sit somewhere between **rules-based automation** (RPA, decision trees) and **GenAI augmentation** (assistants for staff, summarization, coding copilots). Agentic AI is a step change: **systems that can understand a goal, decompose it into tasks, invoke tools and systems, monitor progress, and adapt actions within defined guardrails until the goal is achieved.**[3,5,11,13,14]

Operationally, agentic systems differ from previous generations in four ways:

- **Goal-oriented execution:** Agents are configured against business objectives ("resolve KYC case within policy and SLA"), not just discrete steps.
- **Tool use and environment interaction:** Agents call core banking APIs, workflow engines, document systems, market data, rating engines, and even other agents.[13]
- **Stateful, multi-step reasoning:** Agents maintain context across days or weeks, handling dependencies, waiting for external events, and resuming without human prompts.
- **Policy-constrained autonomy:** Agents operate within codified policies (credit, AML, conduct, capital) and automatically escalate exceptions.

The result is **end-to-end workflows where humans supervise and intervene selectively rather than execute the majority of steps.**

Where Agentic AI Is Already Running End-to-End

Below are representative live use cases in 2025–2026 at Tier-1 institutions and fintechs; each is materially beyond "co-pilot" usage.

Autonomous Credit Remediation

Operational mechanism

- Trigger: Past-due or early-warning signals on retail and SME books.
- The agent:
  - Ingests customer transaction patterns, bureau data, collateral values, and macro indicators.
  - Classifies cases (temporary liquidity vs structural impairment).
  - Proposes remediation paths (payment holidays, restructuring, limit reduction) under pre-approved policy templates.
  - Orchestrates outreach (personalized messages, outbound calls), configures revised payment plans, updates core loan systems, and books accounting entries.
  - Monitors adherence and triggers further actions or escalation.

Workflow change

- Previously: Collections agents manually reviewed accounts, generated letters, negotiated terms, keyed updates into disparate systems—often over several weeks.
- Now: **Agentic system handles 60–80% of cases without human intervention, with staff focusing on high-risk, high-complexity exposures.**

P&L / risk impact

- Lower credit losses via earlier, more consistent remediation and better segmentation.[4]
- Reduced collections headcount and external agency spend.
- Improved expected loss volatility, supporting capital optimization at the portfolio level.

Multi-Step KYC / CDD Resolution

Operational mechanism

- Trigger: KYC case opens due to onboarding, periodic review, or alert.
- The agent:
  - Gathers data from internal systems, external registries, sanctions lists, adverse media, and beneficial ownership databases.[8]
  - Requests missing documents via omni-channel interaction, checks quality, and performs OCR and entity resolution.
  - Applies policy rules and risk-scoring models to reach a decision.
  - Drafts and files case documentation, including rationale and evidence, in an auditable format.

Workflow change

- Previously: Analysts spent 30–90 minutes per case; high variance in completion time and quality.
- Now: **Agent resolves low- and medium-risk cases end-to-end**, analysts handle only exceptions, complex structures, and escalation.

Economic effect

- 30–50% cycle-time reduction and 30–40% unit cost reduction reported where GenAI and agentic orchestration are combined.[14]
- Lower backlog risk and reduced regulatory penalties due to consistent evidence trails and full-population coverage.

Intraday Liquidity Optimization

Operational mechanism

- Continuous agents ingest real-time payments flows, securities settlements, collateral positions, and intraday credit lines.
- They:
  - Forecast intraday liquidity needs using short-horizon models.
  - Automatically reprioritize payments, recommend use of central bank facilities, and move collateral between clearing houses and internal pools.[1]
  - Simulate stress scenarios and propose intraday funding strategies to minimize cost and regulatory buffer breaches.

Workflow change

- Previously: Treasury teams manually monitored positions, relying on heuristics and large buffers.
- Now: Agents propose and in some cases execute intraday movements within pre-set constraints, with human sign-off thresholds.

Balance sheet impact

- Reduction in idle liquidity buffers and intraday overdraft costs.
- Better LCR/NSFR positioning with less conservatism, releasing balance sheet capacity.

Self-Healing IT and Operations

Operational mechanism

- Agents monitor logs, incident tickets, capacity metrics, and user journeys.
- They:
  - Detect anomalies (latency spikes, error codes, abandonment in journeys).
  - Correlate issues across systems.
  - Execute pre-approved remediation playbooks (restart services, clear queues, reroute traffic, trigger fallbacks).
  - File incident records and generate post-mortem summaries.

Workflow change

- Previously: NOC and operations teams triaged alerts, executed scripts, and wrote incident reports.
- Now: Agents clear a material share of standard incidents autonomously, with human teams focusing on novel or systemic failures.

Economic effect

- Higher platform availability, lower operational losses, reduced IT support FTEs, improved regulatory incident reporting quality.

Distinguishing Partial Automation from True Agentic Orchestration

True agentic workflows exhibit four characteristics not seen in legacy automation:

- **End-to-end ownership:** The agent carries a case from trigger to closure, not just a step (e.g., from alert to resolved KYC case, not just document classification).
- **Dynamic decision-making:** The agent can change plan based on new information (e.g., customer response, new alert).
- **Toolchain breadth:** The agent interacts with multiple internal and external systems, not just a single application.
- **Policy-bounded autonomy:** Decisions and actions are constrained by codified limits (e.g., max restructurings, value thresholds), with automatic escalation when exceeded.

Many banks today misclassify enhanced RPA or single-step GenAI as “agents”. Regulators and boards will increasingly expect evidence of **orchestrated, audit-ready autonomy** rather than point solutions.

Figure 1: The Agentic AI Maturity Landscape in Financial Services (2026)

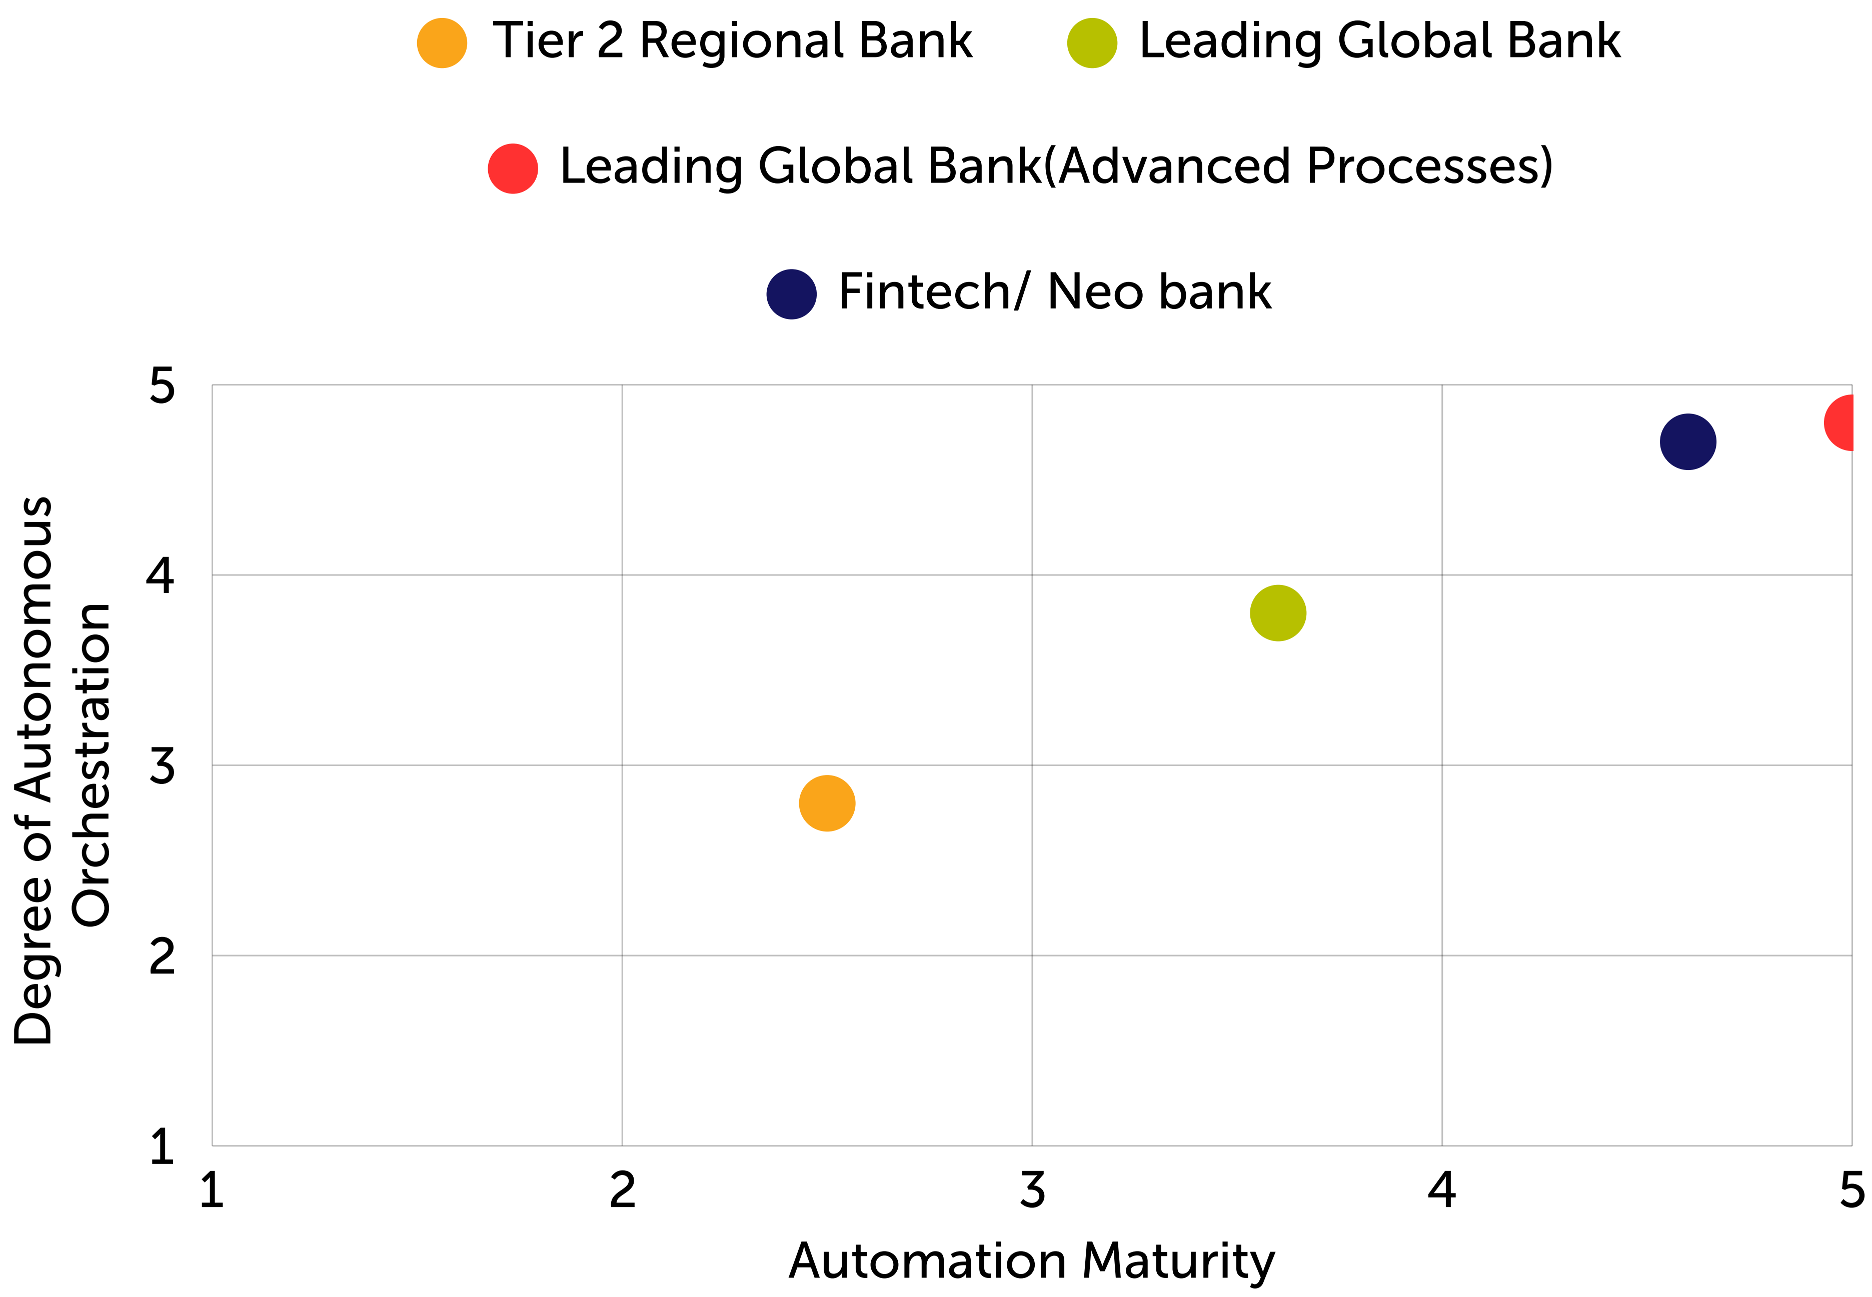

1 - Rules Based Automation; 2 - Advanced Analytics; 3 - GenAI Assistants; 4 - Semi Agentic; 5 - Fully Agentic

2

The Economic Case: Quantifying Value at Stake

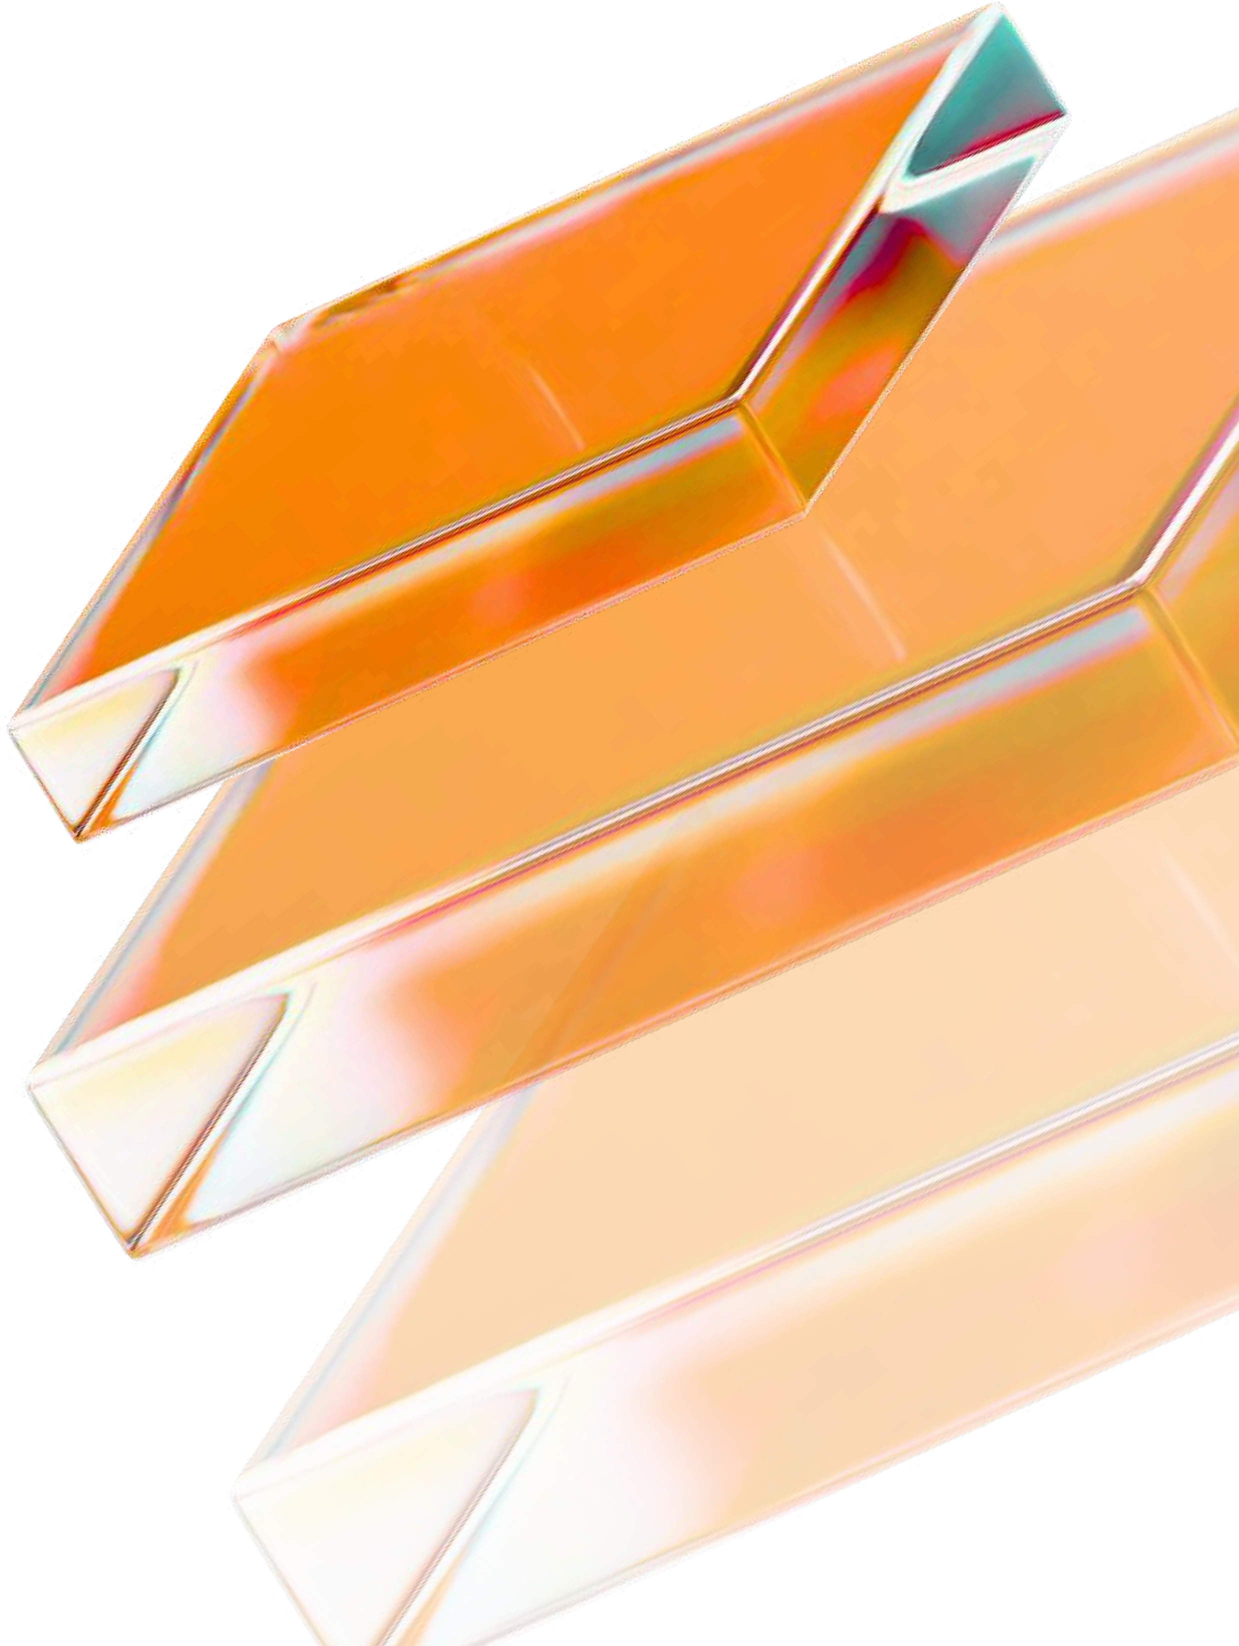

Boards in 2026 are no longer interested in "productivity potential" slides. They are asking three questions:

How does this move the efficiency ratio in the next 12–36 months?

How does this improve ROE and regulatory capital utilization?

What is the downside if we do not match peers' agentic deployments?

35%

Agentic AI increases bank profitability through cost efficiency, risk reduction, and revenue intelligence.

Figure 2: Modeled Profitability Impact of Agentic AI Adoption

Baseline P&L   Agentic P&L

2.1 Cost and Productivity: Middle- and Back-Office

Across financial services, GenAI has already shown 20–40% productivity improvement in core operations and content-heavy tasks.[9,14] Agentic AI amplifies this by:

- Compressing **entire workflows**, not just tasks.
- Enabling **parallelization**: different agents working simultaneously on underwriting, KYC, documentation, and collateral checks.[14]
- Reducing **re-work and exception volumes** through consistent policy application.

Plausible 2026 benchmarks from live deployments and credible projections suggest:

- **Middle/back-office cost per case** in agentified domains down 30–50%.[3]
- **Servicing cost reductions** of ~30% in functions combining GenAI and automation have already been reported by major banks, alongside selective headcount reductions in covered areas.[3]
- End-to-end automation of specific workflows (e.g., KYC for low-risk retail) achieving straight-through processing rates exceeding 70–80%.

Translated to a typical universal bank:

- If 30–40% of middle/back-office FTE-addressable work is agentifiable over 3–5 years, and agents reduce unit cost in those domains by 30–40%, **group-level cost bases can fall 8–12% at constant volume.**
- With revenues flat to slightly up, this alone can **improve cost-to-income ratios by 3–5 percentage points**—a step-change shareholders will notice.

2.2 Operational Loss Avoidance and Risk

Agentic systems running continuous monitoring and remediation yield:

- **Fraud and AML:** Enhanced detection precision and real-time intervention, reducing false positives and missed cases.[7]
- **Operational risk:** Automated reconciliation and anomaly detection in back-office and IT operations reduce incidents and manual errors.[5]
- **Compliance:** Continuous surveillance and dynamic rule application lower the probability of systematic breaches and material fines.[11]

Recent industry analyses estimate that AI in banking could increase aggregate industry profits by approximately 9%, partly through reduced operational losses and risk costs. When agents move from alerting to **acting**, a larger share of this benefit becomes realizable —through[9];

- Lower **expected operational losses** in RCSA and ICAAP models.
- Reduced **regulatory capital add-ons** tied to control failures.
- Faster remediation of findings, reducing remediation program spend.

2.3 Revenue Uplift: From Rules-Based Offers to Agentic Personalization

Traditional cross-sell engines rely on:

- Static segmentation.
- Batch-based product propensity models.
- Limited real-time context (often only balances and basic events).

Agentic revenue systems differ in three ways:

- **Continuous context ingestion:** Transactions, digital interactions, life-event signals, external data, and macro conditions.[2]
- **Goal-driven optimization:** Agents are set explicit objectives such as "maximize risk-adjusted NII and fee income for segment X under customer conduct and suitability constraints".
- **Closed-loop experimentation:** Agents run controlled experiments (messaging, channel, offer structure) and adjust strategies continuously.

Concrete impacts:

- Personalized wealth and retail advisors already demonstrating 10–20% uplift in product penetration and AUM growth, with higher customer satisfaction.[8]
- Proactive credit and deposit management improving margins by optimizing pricing and utilization at a customer and segment level.

From a CFO's perspective:

- A 10–15% uplift in **fee and NII revenue** in target segments (e.g., mass affluent, SME) can add 1–2 percentage points to **group ROE**, especially when combined with a lower cost base.

2.4 Translating to Efficiency Ratio and ROE

Consider a stylized Tier-1 bank:

- Revenue: 100 units.
- Operating costs: 60 units (efficiency ratio 60%).
- Equity: 10 units; net income: 1 unit (ROE 10%).

Scenario over 3–5 years with agentic AI in selected domains:

- Cost reduction on addressable base: 10% (6 units).
- Incremental risk-adjusted revenue: +5 units (through personalization, faster loan throughput, lower leakage).
- New state: revenue 105, costs 54 ⇒ efficiency ratio ~51%.
- If credit quality and capital remain stable, net income could move toward 2–2.5 units (ROE 20–25%) in target portfolios, with group ROE improving 2–5 percentage points depending on mix.

Even if only half of this potential is realized at group level due to constraints, a **2–3 point ROE uplift** is strategically decisive in markets where peers compete for the same capital and M&A optionality.

3

The Capability Chasm: Why Most Banks Are Stuck

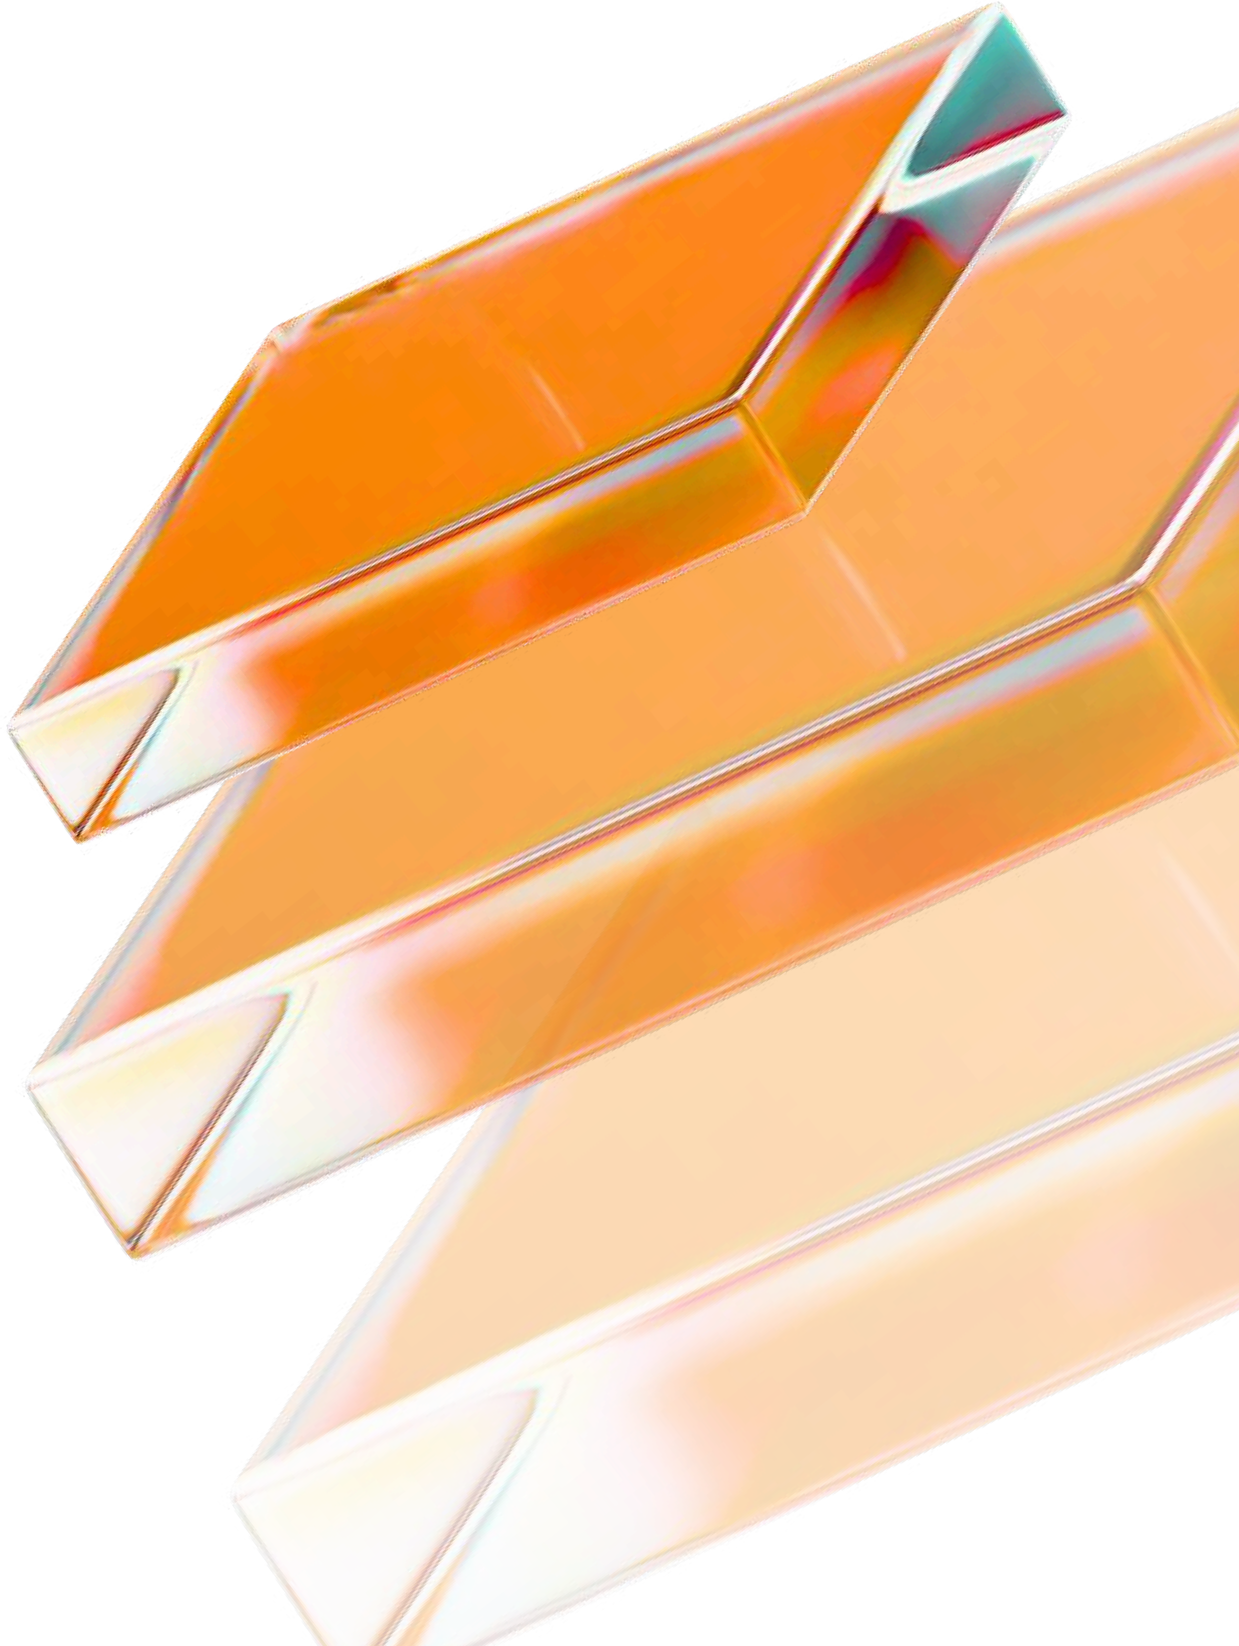

This is the central issue: **the constraint is not "AI models" it is the bank s operating model, architecture, and governance.** Most institutions that assume they can "build agentic capability internally" underestimate the depth and breadth of transformation required.

3.1 Structural Blocker 1 – Legacy Core Systems and Brittle Integrations

Agentic systems require:

- Reliable, performant APIs into **core banking, payments, risk engines, treasury, and document repositories.**
- Event streams that capture **state changes** in near-real-time.
- Idempotent, auditable actions (e.g., book transaction, adjust limit, update KYC record).

Most banks still operate:

- Highly customized cores with **batch-oriented processing,** limiting the ability of agents to act intraday or in real time.
- Fragile integration layers with **point-to-point interfaces,** making orchestration brittle.
- Inconsistent data semantics across product systems, making it hard for agents to compose actions safely.

The result:

- Agents are constrained to the **edges** (e.g., drafting, triage, advisory) rather than executing end-to-end workflows.
- Internal teams struggle to expose **safe, well-governed action APIs** at the speed required.

3.2 Structural Blocker 2 – Absence of Orchestration and Systems-Thinking Talent

Traditional AI staffing in banks has focused on:

- Model development (data scientists).
- MLOps for supervised models.
- Prompt engineering and LLM integration.

Agentic deployments require additional, scarce capabilities:

- **Orchestration architects** who can map end-to-end value streams and define how agents, humans, and systems interact.
- **Agent behavior designers** who encode policies, escalation rules, reward functions, and safe-fail mechanisms.
- **Cross-domain engineers** comfortable with core banking, risk, and workflow engines, not just cloud stacks.

Most banks:

- Have **isolated AI centers of excellence** with limited authority over process redesign.
- Lack deep experience with **multi-agent systems**, tool-use frameworks (e.g., LangGraph-style orchestration), and safety patterns tailored to regulated environments.[13]
- Underestimate the work to move from demos in an innovation lab to **24/7, audited production.**

Hence, internal builds:

- Stall at **pilot scale** or remain trapped in specific silos.
- Fail to prove credible, recurring P&L impact at group level, leading to budget fatigue.

3.3 Structural Blocker 3 – Fragmented Data Ownership and Governance

Agentic AI magnifies both the value and the risk of data:

- Agents need unified views of **customers, exposures, transactions, and documents** to act intelligently.[8]
- They must be constrained by **data minimization, consent, and purpose limitation** rules.

Most banks:

- Have fragmented data ownership across business units with competing priorities.
- Maintain parallel "golden sources" with **inconsistent keys and lineage.**
- Operate data governance as a **compliance function**, not as an integral part of product and agent design.

Consequences:

- Agentic projects either **rebuild narrow, local data pipelines** (limiting scale), or they stall waiting for enterprise data remediation programs.
- Boards and regulators push back when **lineage, explainability, and access control** are not demonstrably in place.

3.4 Structural Blocker 4 – Cultural Resistance to Delegating Autonomy

The agentic model reassigns decision-making:

- **From individual staff** (analysts, operations, relationship managers) to **codified policy and system behavior.**
- Human roles shift from **"doer" to "delegator/overseer".**

Internal resistance emerges when:

- Middle managers fear **loss of headcount and span of control.**
- Risk and compliance teams see autonomy as a **threat to personal accountability.**
- Frontline staff distrust black-box agents.

Without a deliberate change program:

- Agents are constrained to "suggestions", with humans required to click every step—**destroying the economic case.**
- Institutions remain stuck at GenAI augmentation, unable to achieve true operating leverage.

Why Internal Efforts Usually Fail at Scale

Combining the four blockers, internal build-only strategies face common failure modes:

- **Timeline mismatch:** Core modernization, data remediation, and culture change operate on 5–10 year horizons; shareholder and regulatory patience for unclear AI ROI is 2–3 years at best.

- **Fragmented initiatives:** Line-of-business pilots never cohere into a **platform**, resulting in duplicated effort, inconsistent controls, and unscalable architectures.[13]

- **Governance gridlock:** Without credible external benchmarks and blueprints, CROs and regulators default to conservative positions that block autonomy.

The reality in 2026 is that **the capability chasm is a system problem, not a skills problem**. It requires a re-designed operating model: who owns processes, how decisions are encoded, how technology is governed, and how value is measured.

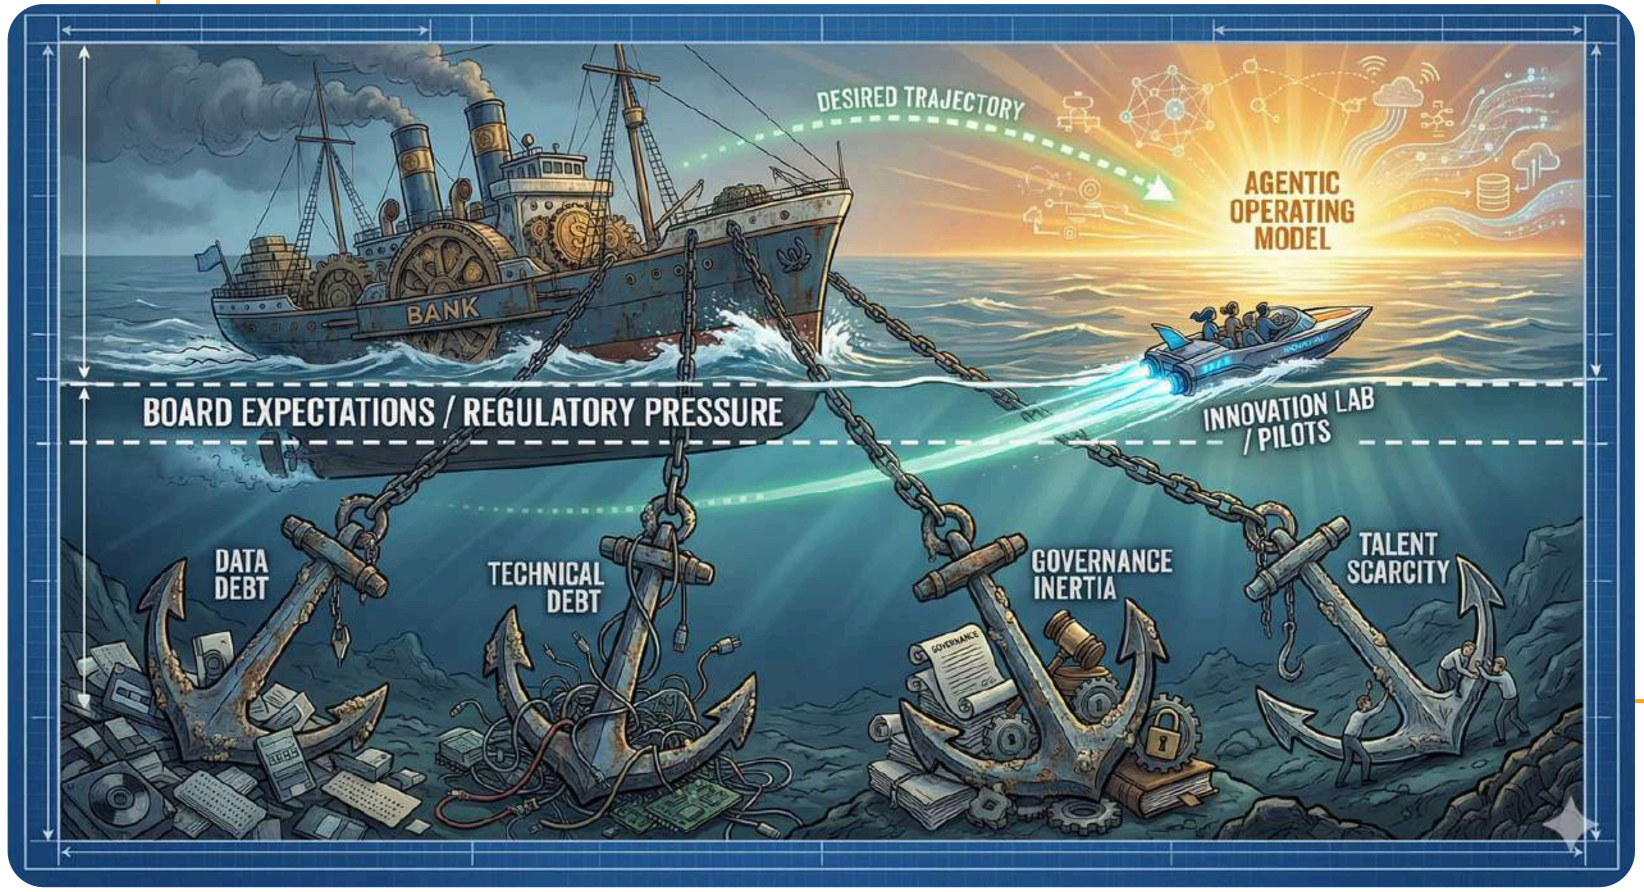

Executive Takeaway: Speedboat-style pilots cannot move the mothership. Without intentionally lifting the four anchors, internal AI programs will remain marginal.

4

Risk Trade-Offs: Adoption Versus Inaction

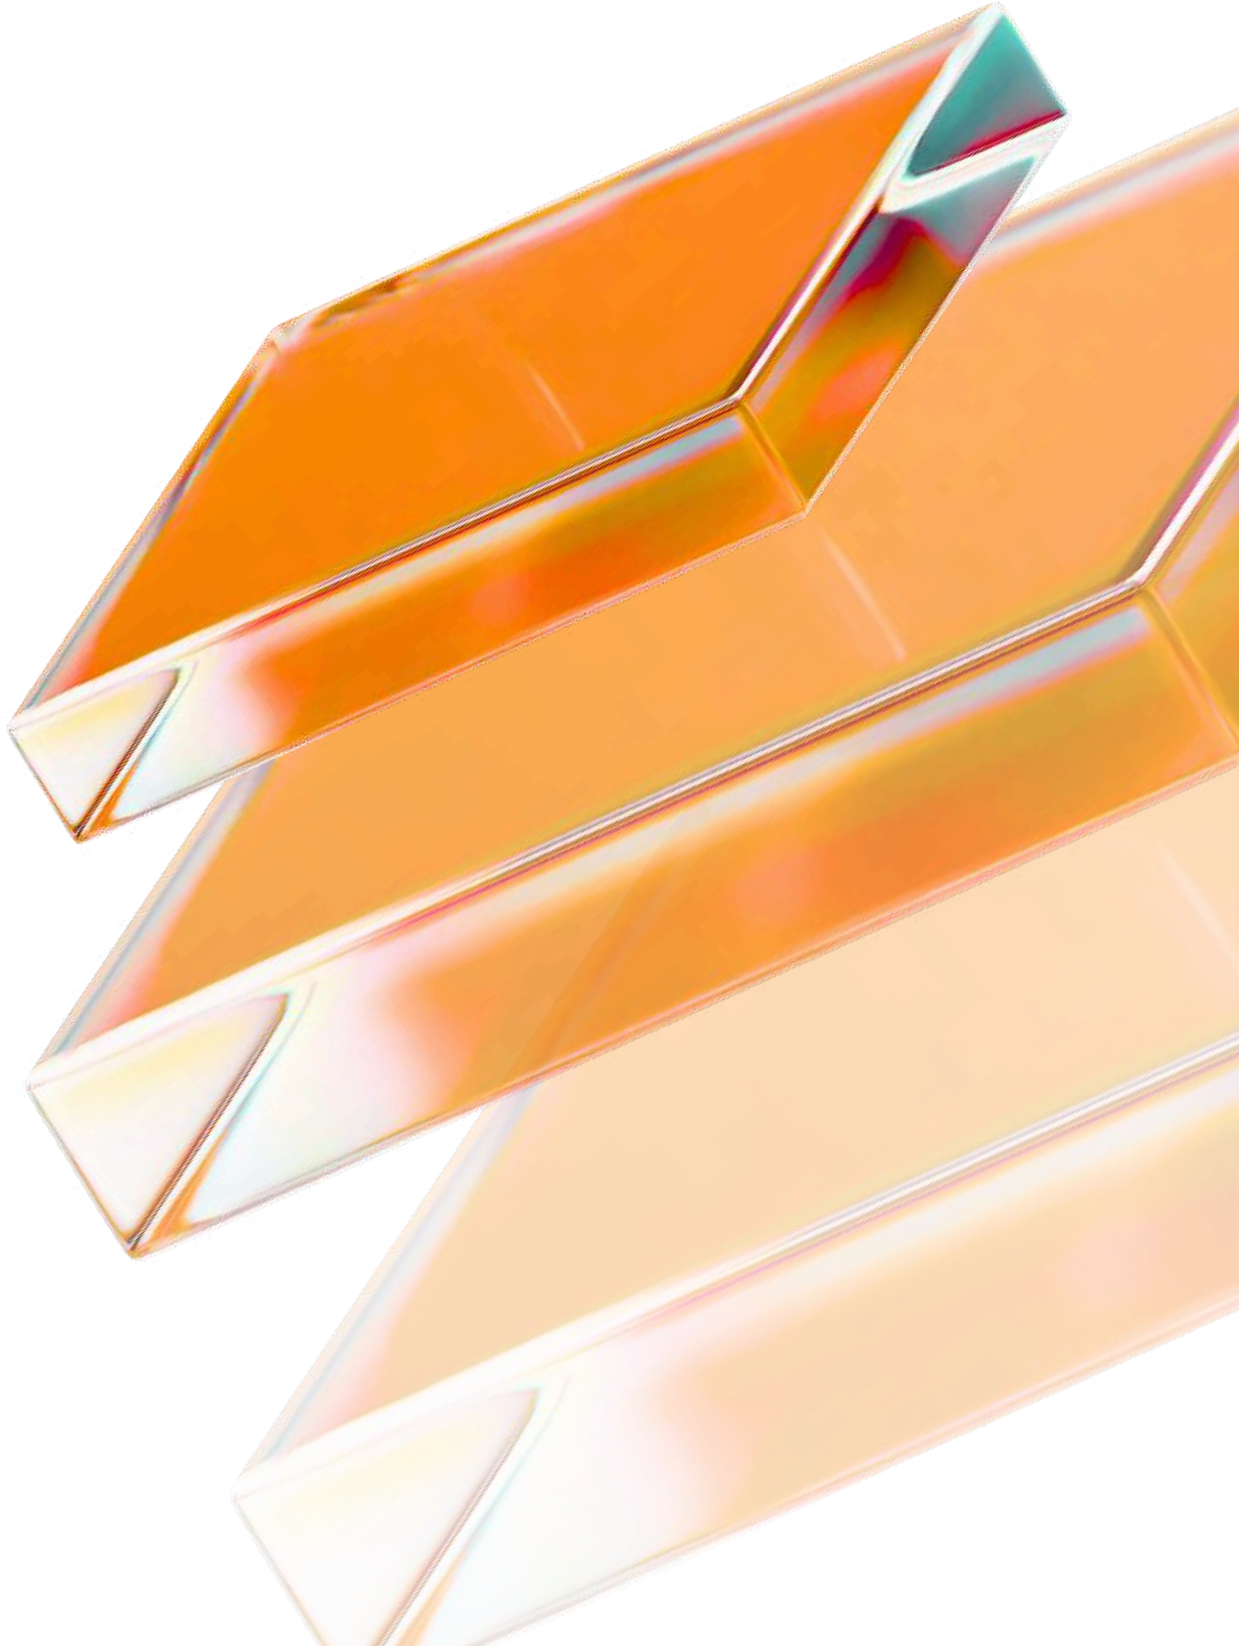

CROs and boards appropriately worry about:

- Hallucinations and **unreliable outputs** from LLMs.
- **Bias** and unfair outcomes in credit, fraud, and pricing.
- **Regulatory scrutiny** over opaque decisioning, especially under emerging AI and data laws.
- Concentration risk in cloud, foundation models, and vendors.

However, leading institutions are demonstrating that **agentic architectures can be designed to be safer than current human-only or rules-based processes**—and that inaction carries greater strategic risk.

4.1 How Leading Banks Are Mitigating AI Risks

Architecture & Model Risk Controls

- **Layered model approach:** Use robust, well-validated predictive models for core risk assessments; use LLMs primarily for **orchestration, explanation, and document interaction,** not for unbounded numerical predictions.[11]

- **Guardrail frameworks:** Implement structured output constraints, policy templates, and hard limits on actions. Agents propose decisions that are checked by rules engines before execution.

- **Tool-only access:** Agents never have raw access to cores; they interact via **narrow, audited APIs** representing pre-approved actions (e.g., “offer plan A/B/C”, “schedule remediation call”, “submit SAR draft”).

Governance and Oversight

- **Explicit accountability:** Assign process-level owners responsible for the end-to-end behavior of agents, with clear RACI matrices.

- **Continuous monitoring:** Implement dashboards and alerts on key risk indicators: decision distributions, exception rates, overrides, fairness metrics, and outcome drift.[11]
- **Human override mechanisms:** For high-impact decisions, agents act in a **propose-then-approve** pattern with configurable thresholds.

Explainability and Auditability

- Agents maintain a **decision log:**
  - Data accessed (with lineage).
  - Models used and their versions.
  - Policy rules invoked.
  - Reasoning summary and rationale.
- Systems produce **human-readable justifications** and can replay decisions during internal and external audits.

These measures align with evolving regulatory expectations around AI explainability, human oversight, and auditable trails.

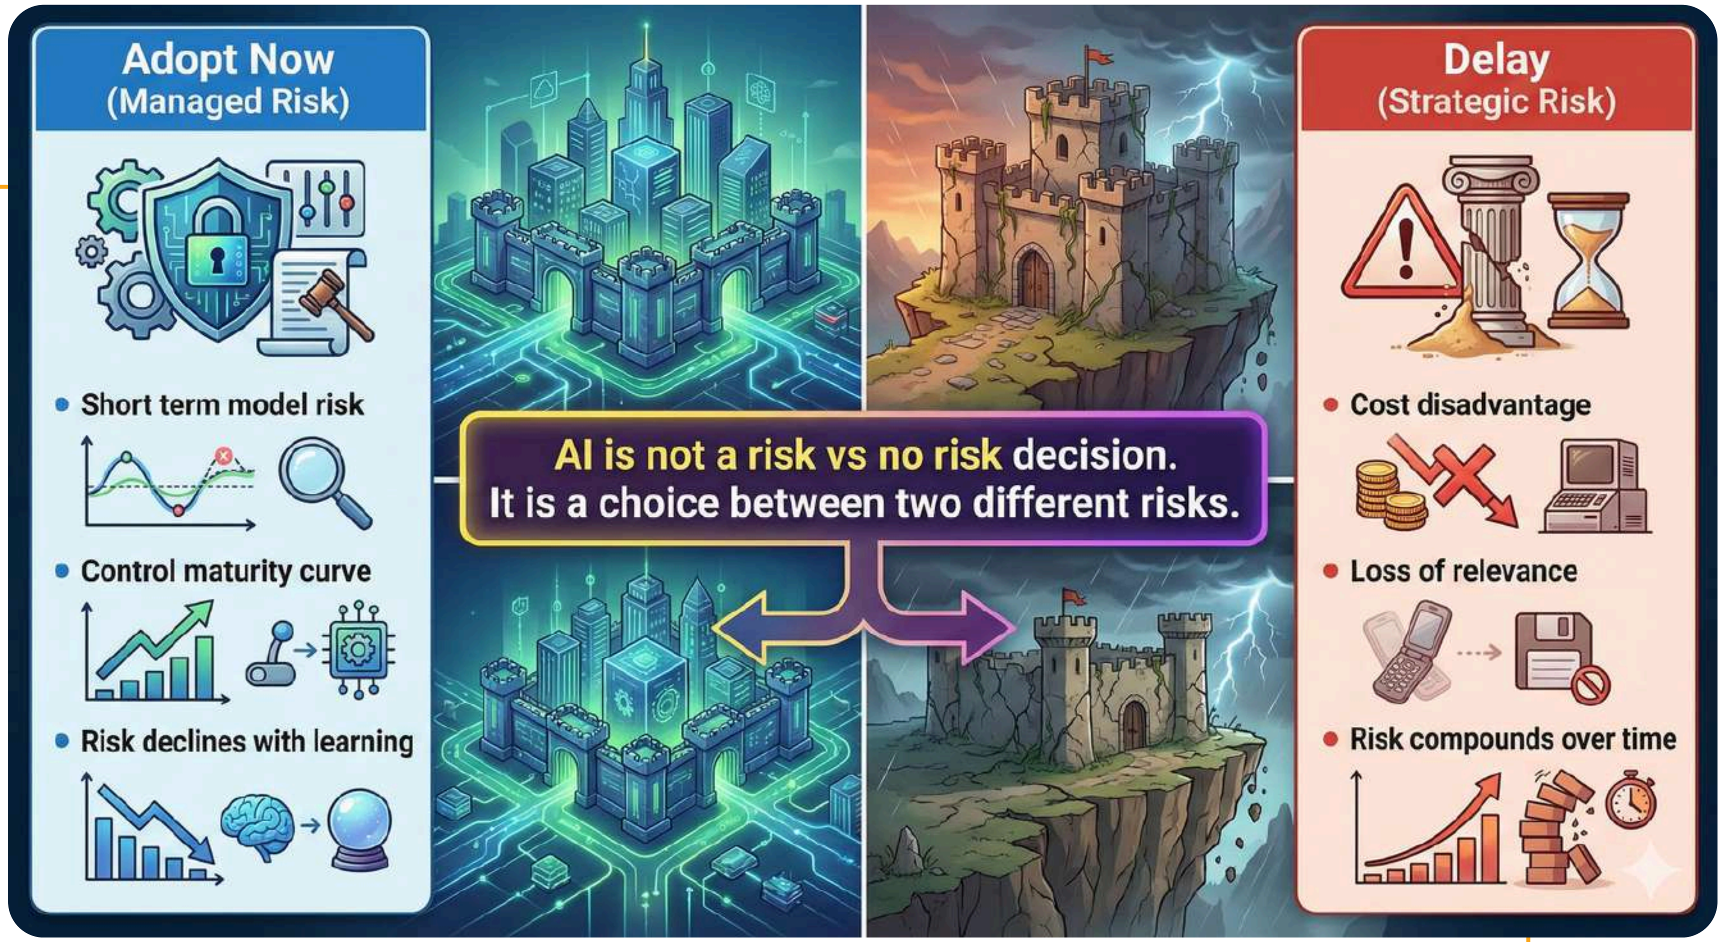

4.2 The Strategic Risk of Delay

The less discussed side of the risk equation is **cost and competitiveness risk:**

- **Structural cost disadvantage:** Banks that cannot reduce unit costs in operations, risk, and IT will have **persistently worse efficiency ratios** (5–10 percentage points in some peer sets) as agentic peers scale.[12]

- **Customer irrelevance:** As fintechs and leading banks offer **agent-based personal finance managers, SME CFO-like agents, and real-time guidance,** traditional banks will appear slow and generic.[10]

- **Regulatory perception:** Supervisors are increasingly aware that advanced AI can reduce risk when properly governed. Institutions that lag may be seen as **less capable of managing complexity,** affecting trust and, over time, capital and liquidity expectations.

In practice, the question for boards is no longer, “Is AI risky?” but rather, “Is not deploying agentic AI at scale a bigger risk to solvency, relevance, and shareholder returns over the next five years?”

4.3 Adoption vs Inaction: A Three-Year Horizon View

Consider:

- Over three years, a leading bank deploying agentic AI in targeted areas can:
  - Remove 8–12% of cost base in addressable functions.
  - Reduce operational losses and compliance findings.
  - Add 5–10% incremental risk-adjusted revenue.
- A laggard maintains current trajectories, perhaps capturing minimal GenAI assistant benefits.

The outcome:

- Relative ROE gaps of several percentage points.
- Diverging price-to-book valuations, as markets reward banks with credible, demonstrable AI-driven operating leverage.
- Reduced strategic options for laggards, which face pressure to merge or exit lines of business.

5

The Only Viable Path Forward

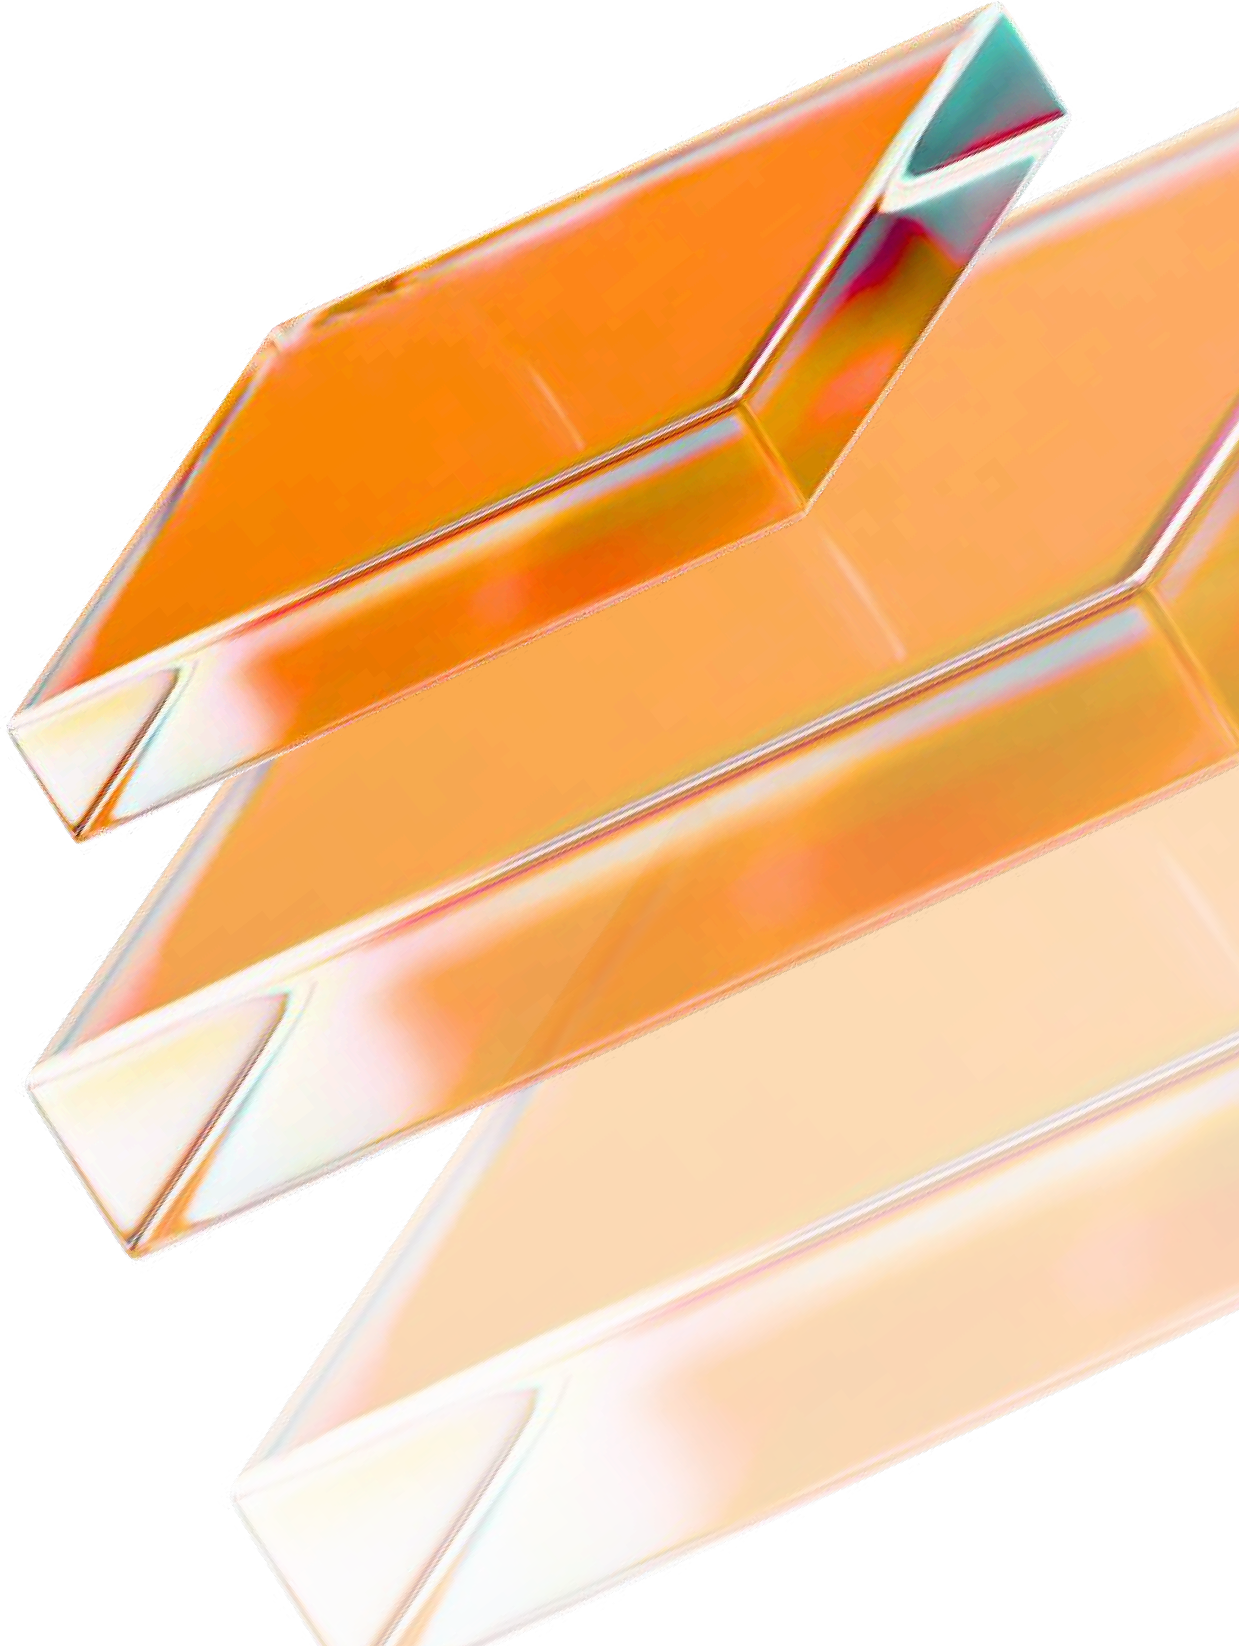

Agentic AI Requires a Different Delivery Model

Achieving the benefits outlined above is not a matter of scaling today's AI labs or buying more models. It requires:

- **Process-first design:** Starting from priority value streams (credit, KYC, treasury, operations) and working backwards to required data, tools, and agents.
- **A dedicated orchestration layer:** An "agentic fabric" that sits between channel/core systems and models, managing workflows, policies, and safety.
- **Integrated data and governance:** A data plane that can support real-time, governed agent access with lineage, consent, and minimization.

Internal organizations optimized for incremental automation are not structurally set up for this. Attempting to retrofit:

- Consumes scarce engineering capacity on plumbing rather than value.
- Prolongs timelines beyond board and market patience.
- Often leads to multiple incompatible platforms within the same bank.

Why External Partnerships Are a Force Multiplier, Not a Dependency

Specialized transformation partners bring:

- **Pre-built agentic patterns and blueprints** for core banking workflows, reducing design and build time.
- Experience with **multi-agent orchestration, safety, and controls** in regulated environments, shortening the learning curve.
- Cross-bank, cross-market benchmarks that inform **credible business cases,** sequencing, and risk appetites.
- Capacity to **run parallel workstreams:** standing up data foundations, orchestration layers, and pilot agents simultaneously—something internal teams rarely achieve alone.

Critically, effective partnerships:

- Operate on **joint execution models,** not outsourcing. Bank teams co-design policies, controls, and operating rhythms.
- Build internal **capability and ownership** over time, ensuring the bank is not locked in to a black-box external platform.

From a CFO/CRO perspective:

- Partnering converts a large, uncertain **capex-like program** with extended timelines into a **phased, value-backed investment** with measurable milestones, risk controls, and up-front clarity on outcomes.
- It limits **architectural missteps** that can become future technical debt.

What 2026 Requires: Commitment, Not Exploration

For a typical Tier-1 or Tier-2 bank, 2026 should be framed internally as:

- **Year 1 of a 3-year agentic operating model transition,** not "Year 4 of AI pilots".
- The year when:
  - The board formally accepts that **GenAI chatbots are table stakes** and no longer a differentiator.
  - A **bank-wide agentic strategy** is approved, tying specific use cases to efficiency ratio, ROE, and risk metrics.
  - A **strategic partnership model** is selected and contracted to accelerate implementation in 2–3 priority domains.
  - Operating model changes (governance, roles, metrics) are defined and agreed across technology, risk, operations, and business lines.

Delaying such a commitment by another 12–24 months effectively concedes structural advantage to more agile peers.

A Decisive, Understated Executive Directive

For boards and executive committees, an appropriate 2026 directive might be stated internally as:

"Over the next three years, this bank will systematically embed agentic AI into its core workflows to achieve step-change improvements in efficiency, ROE, and risk control. We will not attempt to build every capability alone. Instead, we will combine our regulatory and domain expertise with specialized partners to deliver safe, auditable autonomy in carefully selected value streams. 2026 is the year we move from experimentation to execution."

Workstreams	2026: Commit & Architect	2027: Scale & Standardize	2028: Industrialize
Strategy & Governance	Target operating model, risk framework, partner selection	Policy expansion, regulatory alignment	Continuous oversight, optimization
Data & Controls Foundation	Data readiness, access controls, audit foundations	Standardized controls, monitoring	Embedded compliance by design
Agentic Orchestration Layer	Platform selection, reference architecture	Shared services, orchestration patterns	Enterprise scale orchestration
Priority Use Cases (Credit, KYC, Treasury, Ops)	Prioritization, pilots	3–5 workflows live	Majority of value streams agentified
Capability & Culture	Leadership alignment, core team training	Role based enablement	AI native operating culture

The Edge

Agentic AI is not another technology initiative. It is an operating model realignment that determines which banks can sustain competitive economics in the next decade. Institutions that treat it as an extension of their GenAI experimentation phase will, by 2028, find themselves structurally uncompetitive on cost, slow to move capital and risk, and increasingly peripheral to their customers' financial lives. The window to move from curiosity to conviction is narrowing; 2026 is the year in which that decision is made, whether explicitly or by default.

References

[1] Appinventiv, "How Agentic AI Applications in Banking," Jan. 1, 2026. [Online]. Available: https://appinventiv.com/blog/agentic-ai-in-banking/. [Accessed: Jan. 12, 2026].

[2] World Economic Forum, "How Agentic AI will transform financial services," Jun. 2, 2025. [Online]. Available: https://www.weforum.org/stories/2024/12/agentic-ai-financial-services-autonomy-efficiency-and-inclusion/. [Accessed: Dec. 14, 2025].

[3] The UXDA, "AI Gold Rush: Rewriting the CX in Digital Banking," Aug. 4, 2025. [Online]. Available: https://www.theuxda.com/blog/ai-gold-rush-21-digital-banking-ai-case-studies-cx-transformation. [Accessed: Jan. 5, 2026].

[4] Chaione, "How Agentic AI Is Redefining Banking: 10 High-Impact Use Cases," Jul. 13, 2025. [Online]. Available: https://www.chaione.com/blog/agentic-ai-banking-10-high-impact-use-cases. [Accessed: Dec. 22, 2025].

[5] Signity Solutions, "How Agentic AI Is Transforming Financial Services?," Nov. 11, 2025. [Online]. Available: https://www.signitysolutions.com/blog/agentic-ai-in-financial-services. [Accessed: Jan. 10, 2026].

[6] SBS Software, "How generative AI is driving the future of banking," Oct. 9, 2024. [Online]. Available: https://sbs-software.com/insights/generative-ai-future-of-banking/. [Accessed: Dec. 5, 2025].

[7] MIT Technology Review, "Reimagining the future of banking with agentic AI," Sep. 4, 2025. [Online]. Available: https://www.technologyreview.com/2025/09/04/1123023/imagining-the-future-of-banking-with-agentic-ai/. [Accessed: Jan. 2, 2026].

[8] Creatio, "AI Agents in the Finance Industry: Use Cases & Benefits," Sep. 25, 2025. [Online]. Available: https://www.creatio.com/glossary/ai-agents-in-finance. [Accessed: Dec. 18, 2025].

[9] OpenText, "State of AI in Banking | Digital Banking Report." [Online]. Available: https://www.opentext.com/media/report/state-of-ai-in-banking-digital-banking-report-en.pdf. [Accessed: Jan. 8, 2026].

[10] xCube LABS, "Top Agentic AI Use Cases in Banking to Watch in 2025," Aug. 28, 2025. [Online]. Available: https://www.xcubelabs.com/blog/top-agentic-ai-use-cases-in-banking-to-watch-in-2025/. [Accessed: Dec. 29, 2025].

[11] Fingent, "Agentic AI in Financial Services: The Shift from Automation to Autonomy," Sep. 10, 2025. [Online]. Available: https://www.fingent.com/au/blog/agentic-ai-in-financial-services-the-shift-from-automation-to-autonomy/. [Accessed: Jan. 14, 2026].

[12] Databricks, "How Data and AI Are Rewiring Productivity in Banking and Payments," Aug. 11, 2025. [Online]. Available: https://www.databricks.com/blog/powering-growth-how-data-and-ai-are-rewiring-productivity-banking-and-payments. [Accessed: Dec. 7, 2025].

[13] Deloitte, "How banks can supercharge intelligent automation with agentic AI," Dec. 23, 2025. [Online]. Available: https://www.deloitte.com/us/en/insights/industry/financial-services/agentic-ai-banking.html. [Accessed: Jan. 4, 2026].

[14] nCino, "Agentic AI in Banking: How Autonomous AI Is Transforming," Dec. 16, 2025. [Online]. Available: https://www.ncino.com/blog/agentic-ai-banking-revolution-autonomous-intelligence. [Accessed: Dec. 30, 2025].

[15] World Economic Forum, "Artificial Intelligence in Financial Services 2025." [Online]. Available: https://reports.weforum.org/docs/WEF_Artificial_Intelligence_in_Financial_Services_2025.pdf. [Accessed: Jan. 11, 2026].

About Otic Group

Otic Group is a leading AI consulting organisation, driving enterprise transformation through strategic consulting, intelligent automation, and capacity building. Officially endorsed by the Government of Uganda for AI implementation, Otic Group integrates policy, innovation, and execution — helping organizations future-proof their operations and become data-driven enterprises. Our consulting practice aligns AI strategy with corporate vision, delivering measurable business outcomes across key industries.

Visit us at: oticgroup.net

Disclaimer: This research report is for informational purposes only and does not constitute investment, financial, or technology advice. Forward-looking projections and benchmarks are illustrative based on public data as of January 2026; actual results may differ materially. Otic makes no warranties regarding accuracy or completeness and accepts no liability for decisions made in reliance on this analysis. Always consult qualified advisors for institution-specific applicability.